



Week3

Evaluation metrics for binary classifiers

Classification - Accuracy with imbalanced classes:

Dummy classifier

Dummy regressors

Confusion matrices

F-score: Combining precision & recall.

Scores for predictions points

Decision functions (decision_function)

Predicted probability of class membership (predict_proba)

Curves

Precision-Recall

ROC (Receiver Operating characteristic curve)

Multi-class evaluation

Evaluation metrics for binary classifiers

- Represent (Extract and select features) → Train (fit estimator to data) → Evaluate → Refine cycle (feature and model).
- The metric for train and evaluation don't have to be the same.

Classification - Accuracy with imbalanced classes:

- Accuracy = $\frac{\text{\#correct predictions}}{\text{\#total instances}}$.
- It's useless with imbalanced classes, a dummy classifier (for example one that predicts the majority class) can be similar to a model like SVM.

Dummy classifier

- Provides a null metric (accuracy) baseline.
- To be used for model comparison, not for real data problems outcome.
- Serves as a sanity check on performance.
- There are various strategies for the dummy classifier:

- `most_frequent`: predicts the most frequent label in the training set.
- `stratified`: random predictions based on the training set class distribution
- `uniform`: generates predictions uniformly at random.
- `constant` always predicts a constant label provided by the user.
 - Useful with F1 score when the positive class is in the minority.
- If my classifier accuracy is close to the null accuracy baseline, it could indicate:
 - Ineffective, erroneous or missing features.
 - Large class imbalance.
 - Poor choice of kernel or hyperparameters.

Dummy regressors

- Same logic, but the strategy turns into:
 - Mean: Predicts the mean of the training targets
 - Median: predicts the median of the training targets.
 - quantile: predicts a user-provided quantile of the training targets.
 - constant: predicts a constant user provided value.

Confusion matrices

Binary Prediction Outcomes			
<div>True negative</div> <div>True positive</div>	TN	FP	<div>Label 1 = positive class (class of interest)</div> <div>Label 0 = negative class (everything else)</div> <div>TP = true positive</div> <div>FP = false positive (Type I error)</div> <div>TN = true negative</div> <div>FN = false negative (Type II error)</div>
	FN	TP	
	Predicted negative	Predicted positive	

Accuracy: for what fraction of all instances is the classifier's prediction correct (for either positive or negative class)?			
True negative	TN = 400	FP = 7	$\text{Accuracy} = \frac{TN+TP}{TN+TP+FN+FP}$ $= \frac{400+26}{400+26+17+7}$ $= 0.95$
True positive	FN = 17	TP = 26	
	Predicted negative	Predicted positive	
			N = 450

Classification error (1 – Accuracy): for what fraction of all instances is the classifier's prediction **incorrect**?

True negative	TN = 400	FP = 7	
True positive	FN = 17	TP = 26	
	Predicted negative	Predicted positive	N = 450

$$\text{ClassificationError} = \frac{FP + FN}{TN + TP + FN + FP}$$

$$= \frac{7+17}{400+26+17+7}$$

$$= 0.060$$

Recall, or True Positive Rate (TPR): what fraction of all positive instances does the classifier **correctly** identify as positive?

True negative	TN = 400	FP = 7	
True positive	FN = 17	TP = 26	
	Predicted negative	Predicted positive	N = 450

$$\text{Recall} = \frac{TP}{TP + FN}$$

$$= \frac{26}{26+17}$$

$$= 0.60$$

Recall is also known as:

- True Positive Rate (TPR)
- Sensitivity
- Probability of detection

Precision: what fraction of **positive** predictions are correct?

True negative	TN = 400	FP = 7	
True positive	FN = 17	TP = 26	
	Predicted negative	Predicted positive	N = 450

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$= \frac{26}{26+7}$$

$$= 0.79$$

False positive rate (FPR): what fraction of all negative instances does the classifier **incorrectly** identify as positive?

True negative	TN = 400	FP = 7	
True positive	FN = 17	TP = 26	
	Predicted negative	Predicted positive	N = 450

$$\text{FPR} = \frac{FP}{TN + FP}$$

$$= \frac{7}{400+7}$$

$$= 0.02$$

False Positive Rate is also known as:

- Specificity

- There is often a tradeoff between precision & recall:
 - Recall oriented ML tasks:
 - Search and information extraction in legal discovery
 - Tumour detection
 - Often paired with a human expert to filter out false positives.
 - Precision oriented ML tasks:
 - Document classification.
 - Search engine ranking, query suggestion.
 - Many customer-facing tasks

F-score: Combining precision & recall.

- General evaluation metric: **F-score**

$$F_{\beta} = (1 + \beta^2) \frac{\text{Precision} \cdot \text{Recall}}{(\beta^2 \text{Precision}) + \text{Recall}} = \frac{(1 + \beta^2) \cdot TP}{(1 + \beta^2) \cdot TP + \beta \cdot FN + FP}$$



β allows adjustment of the metric to control emphasis on recall vs precision:

-Precision orientad: $\beta = 0.5$ (false positives hurt performance more than false negatives)

-Recall orientad: $\beta = 2$ (false negatives hurt performance more than false positives)

- **F1 score** - Harmonic mean of precision and recall ($\beta = 1$):

$$F_1 = 2 \cdot \frac{Precision \cdot Recall}{Precision + Recall} = \frac{2 \cdot TP}{2 \cdot TP + FN + FP}$$



In sklearn: **Classification report** provides all this metrics. **Support** is the # instances in the test set that have true label.

Scores for predictions points

Decision functions (decision_function)

- Each classifier score value per test point indicates how confidently the classifier predicts the positive class (large-magnitude positive values) or the negative class (large-magnitude negative values).
- Choosing a fixed decision threshold gives a classification rule.
- By sweeping the decision threshold through the entire range of possible score values, we get a series of classification outcomes that form a curve.

Predicted probability of class membership (predict_proba)

- Typical rule: choose most likely class
 - e.g class 1 if threshold > 0.50.
- Adjusting threshold affects predictions of classifier.
- Higher threshold results in a more conservative classifier

- e.g. only predict Class 1 if estimated probability of class 1 is above 70%
- This increases precision. Doesn't predict class 1 as often, but when it does, it gets high proportion of class 1 instances correct.
- Not all models provide realistic probability estimates

Curves

Precision-Recall

X-axis: Precision

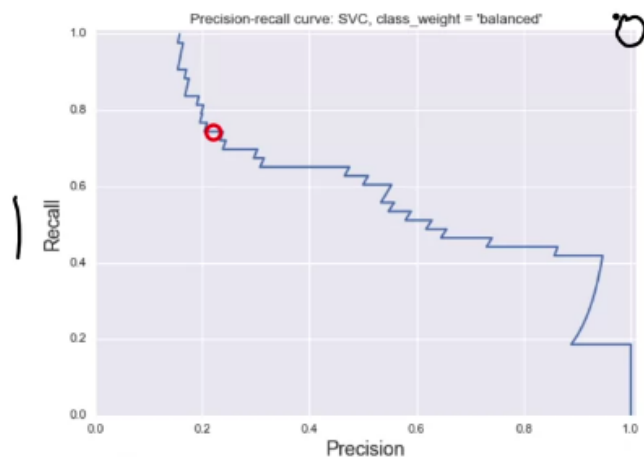
Y-axis: Recall

Top right corner:

- The “ideal” point
- Precision = 1.0
- Recall = 1.0

“Steepness” of P-R curves is important:

- Maximize precision
- while maximizing recall



ROC (Receiver Operating characteristic curve)

X-axis: False Positive Rate

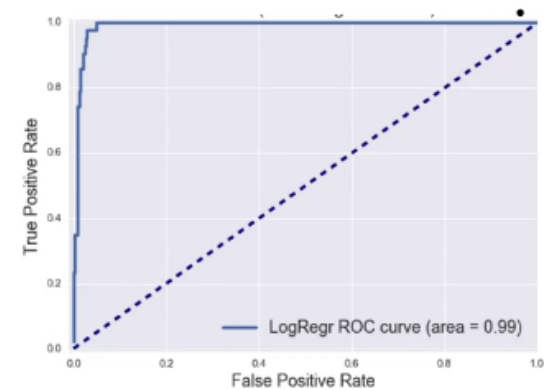
Y-axis: True Positive Rate

Top left corner:

- The “ideal” point
- False positive rate of zero
- True positive rate of one

“Steepness” of ROC curves is important:

- Maximize the true positive rate
- while minimizing the false positive rate



- ROC is the graph resulting from plotting **True Positive Rate(sensitivity or recall) vs. False Positive Rate(1 – specificity)**.
- We can calculate the Area Under the Curve to summarise a classifier performance.
- Advantages:


- Gives a single number.
- Does not require specifying a decision threshold.
- Disadvantages:
 - As with other single number metrics, AUC loses information (shape of the curve, tradeoffs).

Multi-class evaluation

Extension of binary case

- Multi-class evaluation is an extension of the binary case.
 - A collection of true vs predicted binary outcomes, one per class.
 - Confusion matrices are especially useful
 - Classification report
- Overall evaluation metrics are averages across classes
 - But there are different ways to average multi-class results: we will cover these shortly.
 - The support (number of instances) for each class is important to consider, e.g. in case of imbalanced classes
- Multi-label classification: each instance can have multiple labels (not covered here)

Multi-class confusion matrix



Micro vs Macro average

Macro average recall

- Each class has equal weight (doesn't matter it's an imbalanced dataset):
1. Compute the metric (like recall) in each class.
 2. Average resulting metrics across classes.

Micro average recall

- Each instance has equal weight.
 - Largest classes have more influence.
1. Aggregate outcomes across all classes.
 2. Compute metric with aggregate outcomes.
- If the classes have about the same number of instances, macro- and micro-average will be about the same.
 - If some classes are much larger (more instances) than others, and you want to:
 - Weight your metric toward the largest ones, use micro-averaging.
 - Weight your metric toward the smallest ones, use macro-averaging.

- If the micro-average is much lower than the macro-average then examine the larger classes for poor metric performance.
- If the macro-average is much lower than the micro-average then examine the smaller classes for poor metric performance.

Regression evaluation

- Typically R^2 (computes how well future instances will be predicted) is enough.
 - $R^2 = 1$ is a perfect model.
 - $R^2 = 0$ is a model that outputs a constant value regardless of input.
 - R^2 can be negative for bad models.
- Alternative:
 - Mean square error (squared differences of target and predicted).
 - Mean absolute error (abs differences of target and predicted).
 - Median absolute error (robust to outliers).
- There are also dummy regressors to establish a baseline. In the case of a simple regression it can be the mean, median, quantile, custom constant , etc.

Model selection

Optimizing classifiers for different evaluation metrics

- Train/Test on same data
 - Single metric
 - Typically overfits.
 - Can serve as sanity check: low accuracy may indicate a model/implementation problem.
- Single train/test split
 - Single metric.

- Speed and simple.
- Lack of variance information
- K-fold cross-val
 - K train-test split
 - Average metric over all splits.
 - Can be combined with grid search.
 - It can lead to subtle overfitting/optimistic generalisation
- **So, ideally, use 3 sets.**
 - Training (model building)
 - Validation (model selection)
 - Test (model evaluation).
- So, in practice:
 - Create an initial training/test split.
 - Do cross-validation on the training data for model/parameter selection.
 - Save the held-out test set for final model evaluation

Conclusions

- Accuracy is often not the right evaluation metric for many real-world machine learning tasks:
 - False positives and false negatives may need to be treated very differently.
 - Make sure you understand the needs of your application and choose an evaluation metric that matches your application, user, or business goals.
- Examples of additional evaluation methods include:
 - Learning curve: How much does accuracy (or other metric) change as a function of the amount of training data?
 - Sensitivity analysis: How much does accuracy (or other metric) change as a function of key learning parameter values?

